



# 计算机系统基础11



## Exceptional Control Flow: Exceptions & Processes

# Today


- **Exceptional Control Flow**
- **Exceptions**  
异常
- **Processes**

# Control Flow

## ■ Processors do only one thing:

- From startup to shutdown, a CPU simply reads and executes (interprets) a sequence of instructions, one at a time
- This sequence is the CPU's *control flow* (or *flow of control*)

### *Physical control flow*



# Altering the Control Flow

## ■ Up to now: two mechanisms for changing control flow:

- Jumps and branches
- Call and return

React to changes in *program state*

## ■ Insufficient for a useful system:

Difficult to react to changes in *system state*

- Data arrives from a disk or a network adapter
- Instruction divides by zero
- User hits Ctrl-C at the keyboard
- System timer expires

## ■ System needs mechanisms for “exceptional control flow”

# Exceptional Control Flow

- **Exists at all levels of a computer system**
- **Low level mechanisms**
  - 1. **Exceptions**
    - Change in control flow in response to a system event (i.e., change in system state)
    - Implemented using combination of hardware and OS software
- **Higher level mechanisms**
  - 2. **Process context switch**  
上下文
    - Implemented by OS software and hardware timer
  - 3. **Signals**
    - Implemented by OS software
  - 4. **Nonlocal jumps**: `setjmp()` and `longjmp()`
    - Implemented by C runtime library

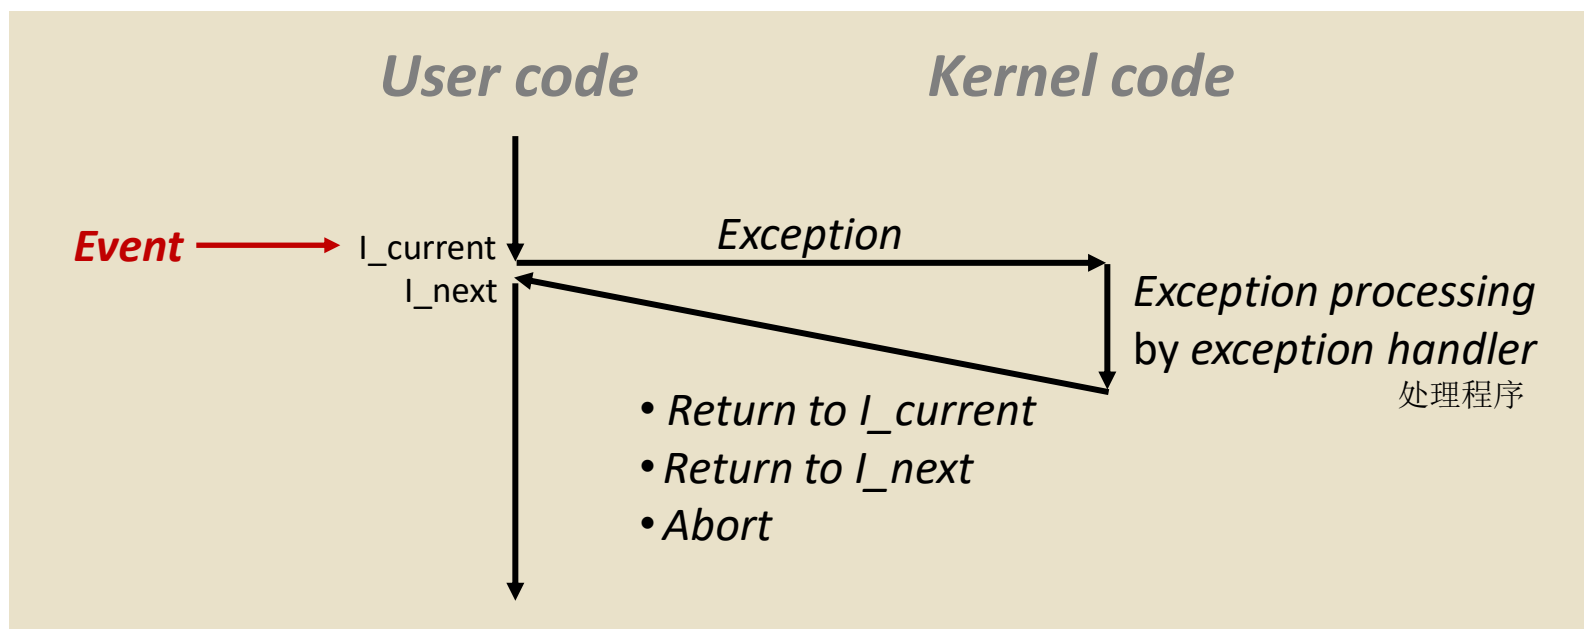
see fig 8-43

# Today

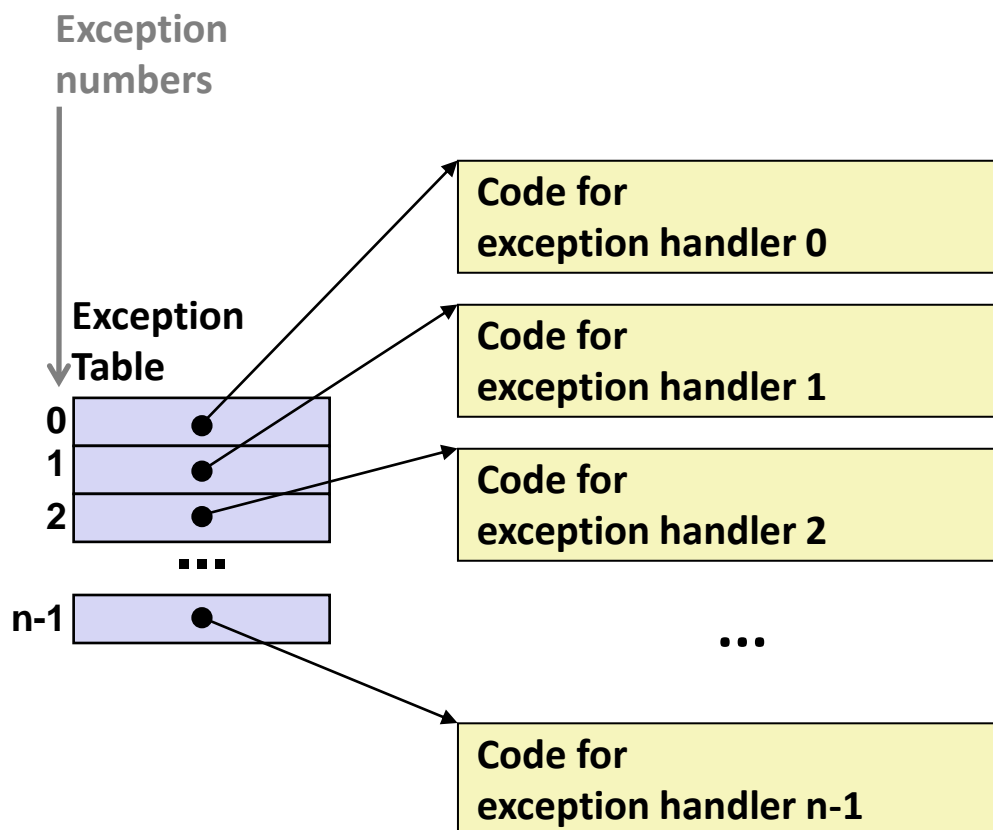
- Exceptional Control Flow
- **Exceptions**
- Processes

# Exceptions

- An **exception** is a transfer of control to the OS *kernel* in response to some *event* (i.e., change in processor state)
  - Kernel is the memory-resident part of the OS
  - Examples of events: Divide by 0, arithmetic overflow, page fault, I/O request completes, typing Ctrl-C



# Exception Tables

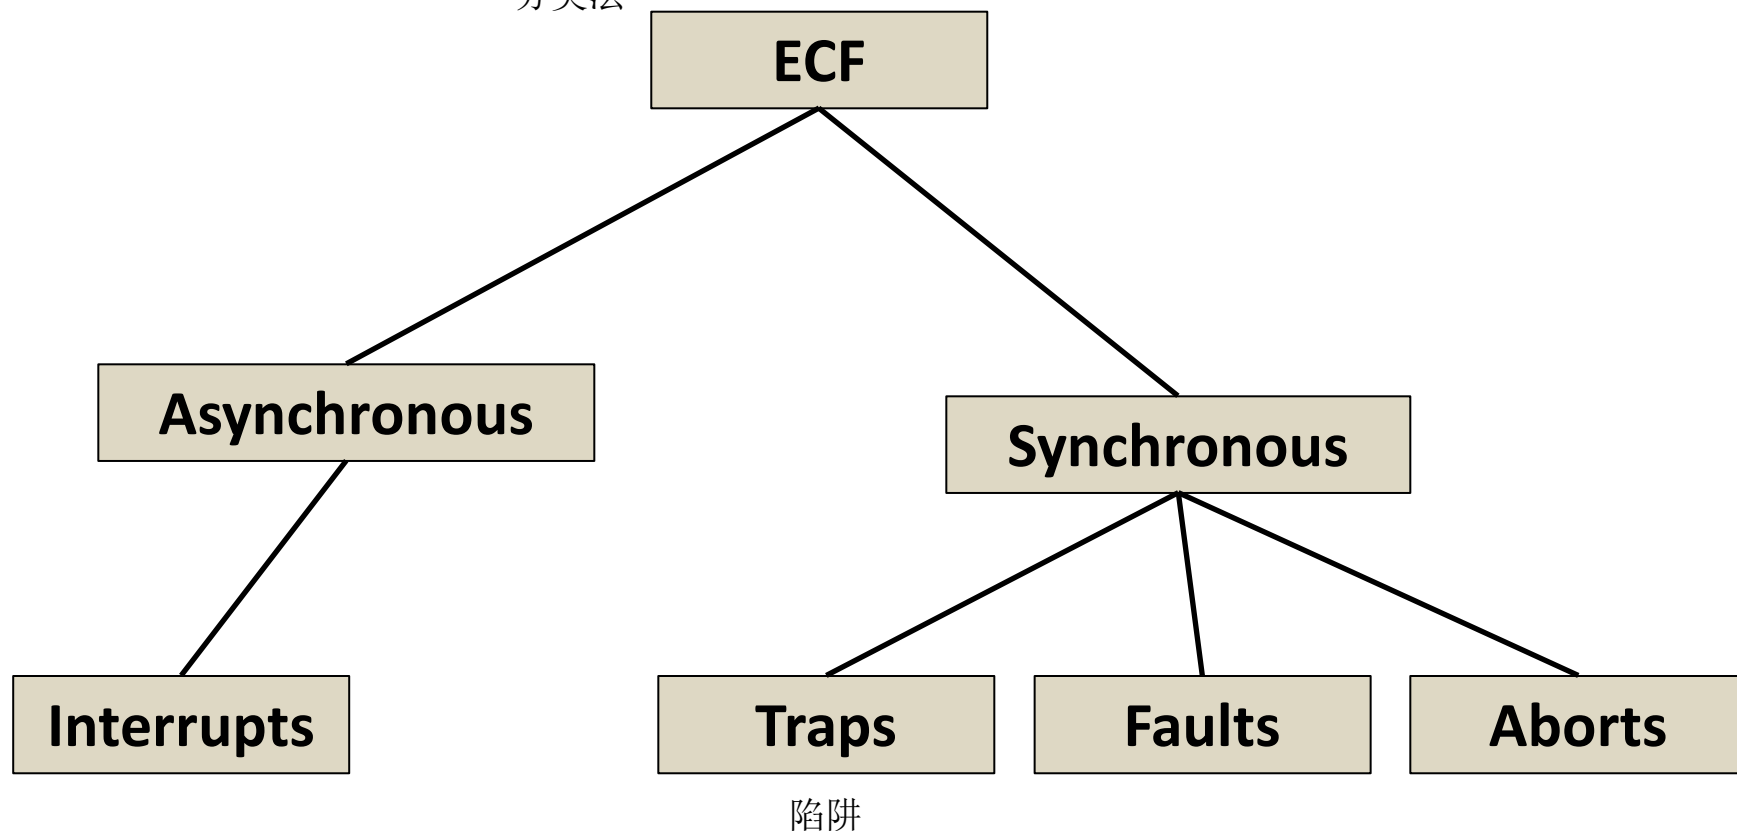


- Each type of event has a unique exception number  $k$  (0-255)
- $k$  = index into exception table (a.k.a. interrupt vector)
- Handler  $k$  is called each time exception  $k$  occurs



# (partial) Taxonomy

分类法



# Asynchronous Exceptions (Interrupts)

## ■ Caused by events external to the processor

- Indicated by setting the processor's *interrupt pin*  
引脚
- Handler returns to “next” instruction

## ■ Examples:

- Timer interrupt
  - Every few ms, an external timer chip triggers an interrupt
  - Used by the kernel to **take back control from user programs**
- I/O interrupt from external device
  - Hitting Ctrl-C at the keyboard
  - Arrival of a packet from a network
  - Arrival of data from a disk

# Synchronous Exceptions

- **Caused by events that occur as a result of executing an instruction:**
  - ***Traps***
    - Intentional, set program up to “trip the trap” and do something
    - Examples: ***system calls***, breakpoint traps, special instructions
    - Returns control to “next” instruction
  - ***Faults***
    - Unintentional but possibly recoverable
    - Examples: page faults (recoverable), protection faults (unrecoverable), floating point exceptions
    - Either **re-executes** faulting (“current”) **instruction** or aborts
  - ***Aborts***
    - Unintentional and unrecoverable
    - Examples: illegal instruction, parity error, machine check
    - Aborts current program

# System Calls

- Each x86-64 system call has a unique ID number
- Examples:

<i>Number</i>	<i>Name</i>	<i>Description</i>
0	read	Read file
1	write	Write file
2	open	Open file
3	close	Close file
4	stat	Get info about file
57	fork	Create process
59	execve	Execute a program
60	_exit	Terminate process
62	kill	Send signal to process

# System Call Example: Opening File

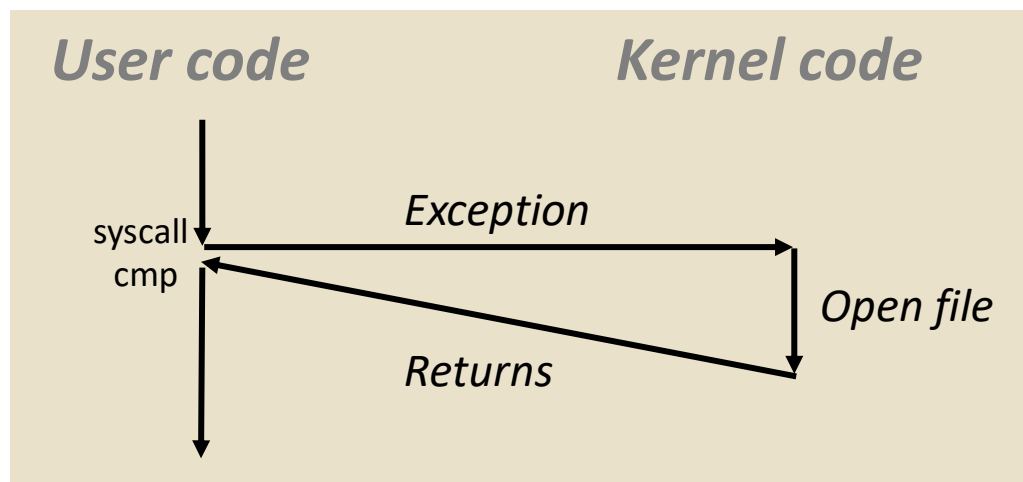
- User calls: `open(filename, options)`
- Calls `__open` function, which invokes system call instruction `syscall`

0000000000e5d70 <\_\_open>:

```

...
e5d79: b8 02 00 00 00  mov  $0x2,%eax           # open is syscall #2
e5d7e: 0f 05             syscall           # Return value in %rax
e5d80: 48 3d 01 f0 ff ff  cmp  $0xfffffffffff001,%rax
...
e5dfa: c3               retq

```



- `%rax` contains syscall number
- Other arguments in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`
- Return value in `%rax`
- Negative value is an error corresponding to negative `errno`

# System Call

- User calls: `open (f`
- Calls `__open` function

```
0000000000e5d70 <__open
...
e5d79: b8 02 00 00 00    m
e5d7e: 0f 05             sy
e5d80: 48 3d 01 f0 ff ff  cr
...
e5dfa: c3               ret
```


Almost like a function call

- Transfer of control
- On return, executes next instruction
- Passes arguments using calling convention
- Gets result in `%rax`

One Important exception!

- Executed by Kernel
- Different set of privileges
- And other differences:
  - E.g., “address” of “function” is in `%rax`
  - Uses `errno`
  - Etc.

*User code*



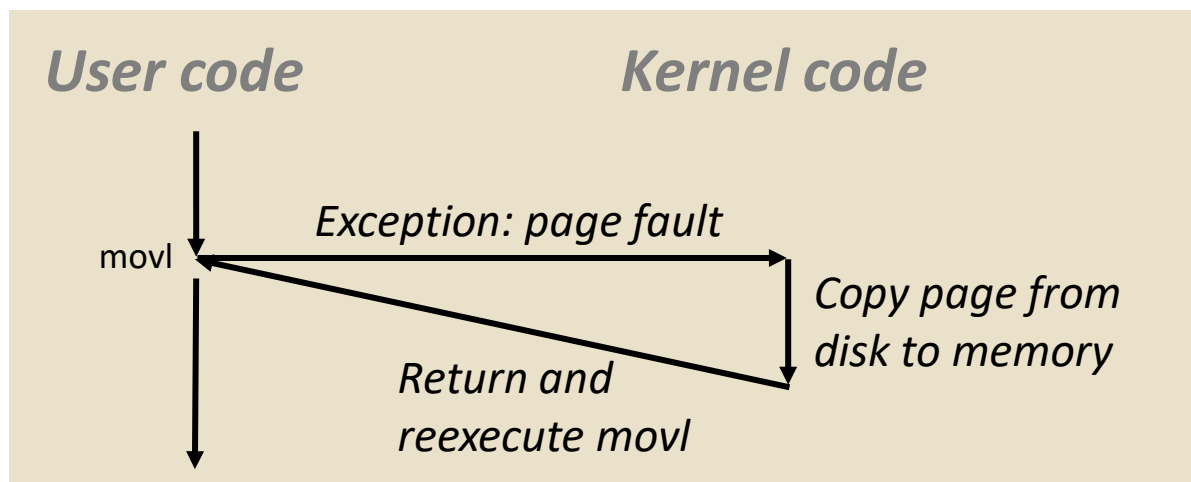
- Return value in `%rax`
- Negative value is an error corresponding to negative `errno`

# Fault Example: Page Fault

- User writes to memory location
- That portion (page) of user's memory is currently on disk

```
int a[1000];
main ()
{
    a[500] = 13;
}
```

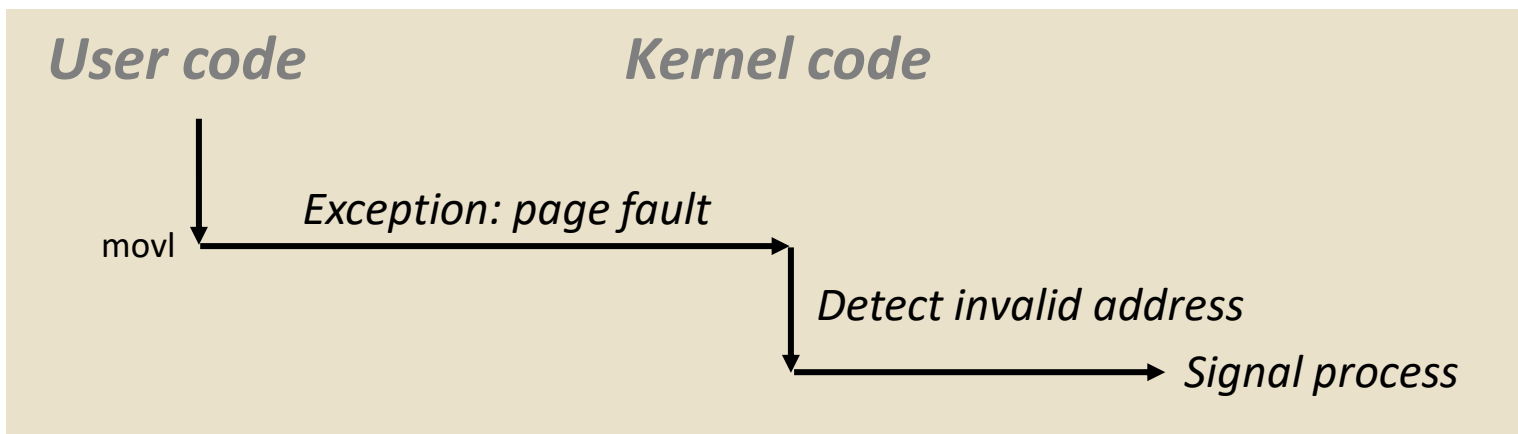
80483b7:	c7 05 10 9d 04 08 0d	movl	\$0xd,0x8049d10
----------	----------------------	------	-----------------



# Fault Example: Invalid Memory Reference

```
int a[1000];  
main ()  
{  
    a[5000] = 13;  
}
```

80483b7:	c7 05 60 e3 04 08 0d	movl	\$0xd,0x804e360
----------	----------------------	------	-----------------



- Sends **SIGSEGV** signal to user process
- User process exits with “segmentation fault”




# Today


- Exceptional Control Flow
- Exceptions
- **Processes**

# Processes

- **Definition:** A *process* is an instance of a running program.
  - One of the most profound ideas in computer science
  - Not the same as “program” or “processor”
  
- **Process provides each program with two key abstractions:**
  - *Logical control flow*
    - Each program seems to have exclusive use of the CPU
    - Provided by kernel mechanism called *context switching*
  - *Private address space*
    - Each program seems to have exclusive use of main memory.
    - Provided by kernel mechanism called *virtual memory*

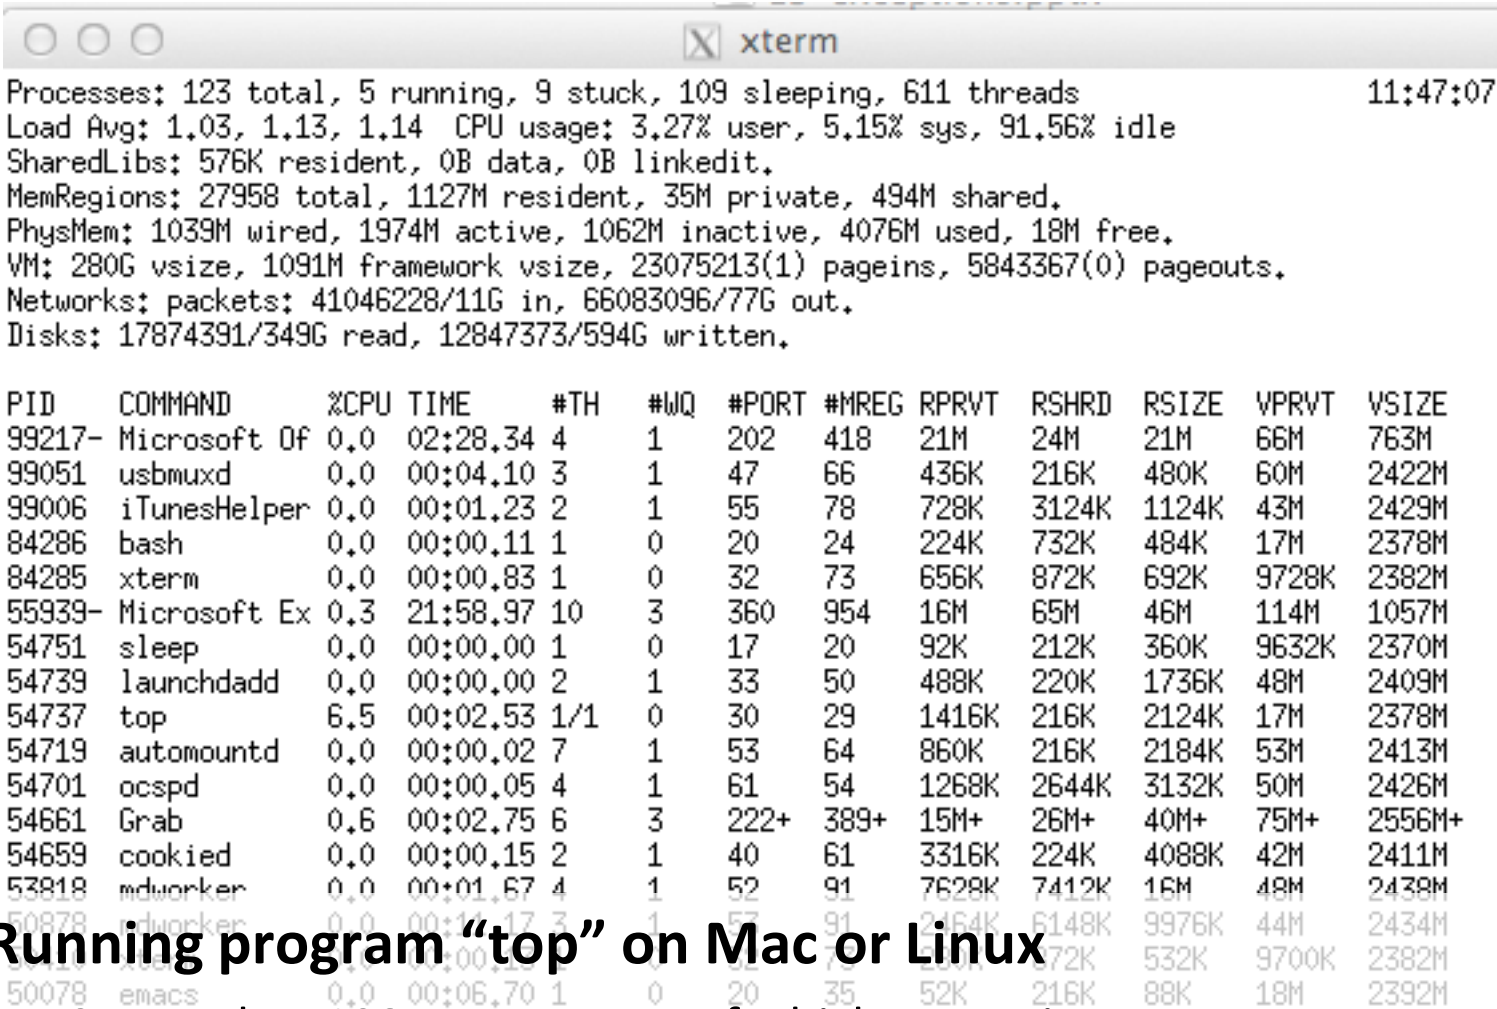


# Multiprocessing



- **Computer runs many processes simultaneously**
  - Applications for one or more users
    - Web browsers, email clients, editors, ...
  - Background tasks
    - Monitoring network & I/O devices

# Multiprocessing Example



```

Processes: 123 total, 5 running, 9 stuck, 109 sleeping, 611 threads
Load Avg: 1.03, 1.13, 1.14  CPU usage: 3.27% user, 5.15% sys, 91.56% idle
SharedLibs: 576K resident, 0B data, 0B linkedit.
MemRegions: 27958 total, 1127M resident, 35M private, 494M shared.
PhysMem: 1039M wired, 1974M active, 1062M inactive, 4076M used, 18M free.
VM: 280G vsize, 1091M framework vsize, 23075213(1) pageins, 5843367(0) pageouts.
Networks: packets: 41046228/11G in, 66083096/77G out.
Disks: 17874391/349G read, 12847373/594G written.

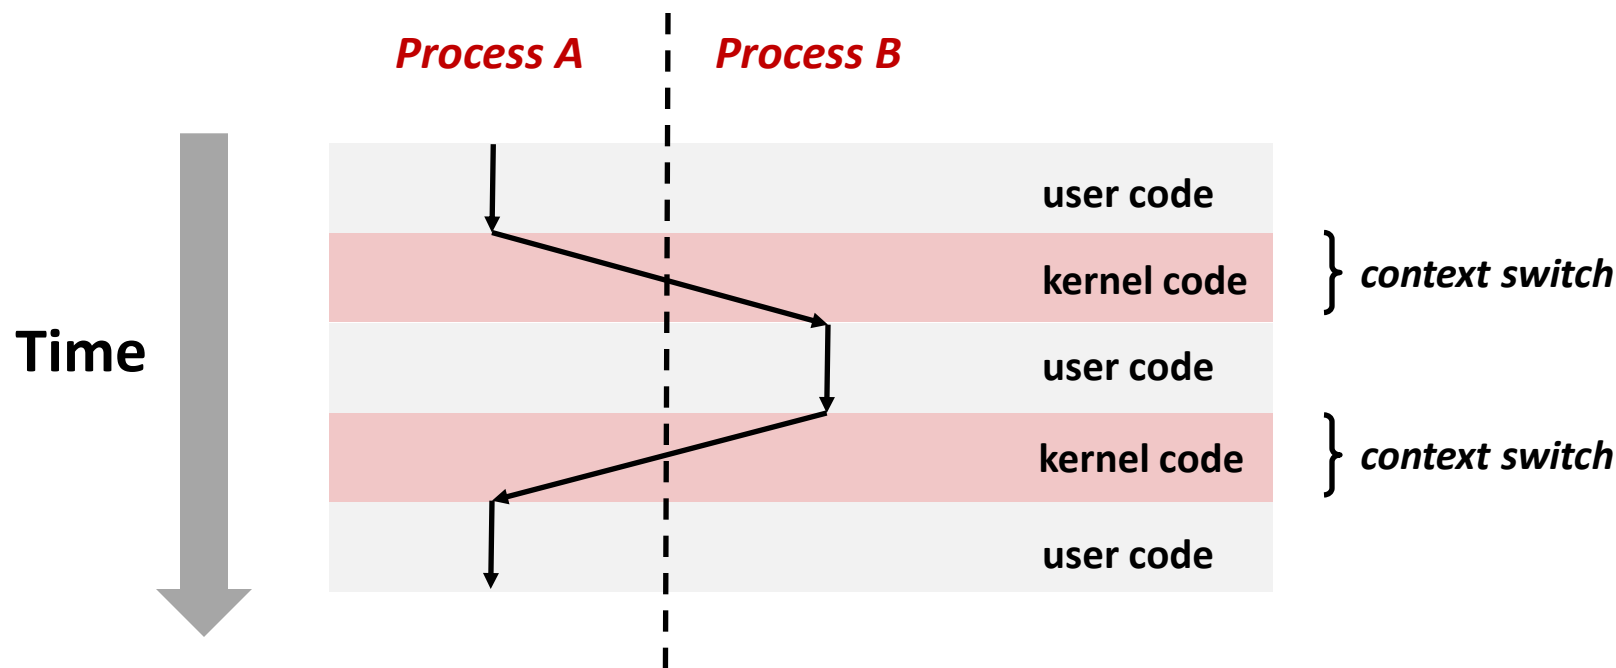
PID    COMMAND      %CPU TIME    #TH   #WQ   #PORT #MREG RPRVT  RSHRD  RSIZE  VPRVT  VSIZE
99217-  Microsoft Of 0.0 02:28.34 4     1    202  418   21M   24M   21M   66M   763M
99051   usbmuxd      0.0 00:04.10 3     1     47   66    436K  216K  480K  60M   2422M
99006   iTunesHelper 0.0 00:01.23 2     1     55   78    728K  3124K 1124K 43M   2429M
84286   bash         0.0 00:00.11 1     0     20   24    224K  732K  484K  17M   2378M
84285   xterm        0.0 00:00.83 1     0     32   73    656K  872K  692K  9728K 2382M
55939-  Microsoft Ex 0.3 21:58.97 10    3    360  954   16M   65M   46M   114M  1057M
54751   sleep        0.0 00:00.00 1     0     17   20    92K   212K  360K  9632K 2370M
54739   launchdadd   0.0 00:00.00 2     1     33   50    488K  220K  1736K 48M   2409M
54737   top          6.5 00:02.53 1/1    0     30   29   1416K  216K  2124K 17M   2378M
54719   automountd   0.0 00:00.02 7     1     53   64    860K  216K  2184K 53M   2413M
54701   ocspd        0.0 00:00.05 4     1     61   54   1268K  2644K 3132K 50M   2426M
54661   Grab         0.6 00:02.75 6     3    222+ 389+  15M+  26M+  40M+  75M+  2556M+
54659   cookied      0.0 00:00.15 2     1     40   61   3316K  224K  4088K 42M   2411M
53818   mdworker     0.0 00:01.67 4     1     52   91   7628K  7412K  16M   48M   2438M
50878   mdworker     0.0 00:01.17 3     1     57   91   2164K  6148K  9976K 44M   2434M
50078   emacs        0.0 00:06.70 1     0     20   35    52K   216K  88K   18M   2392M

```

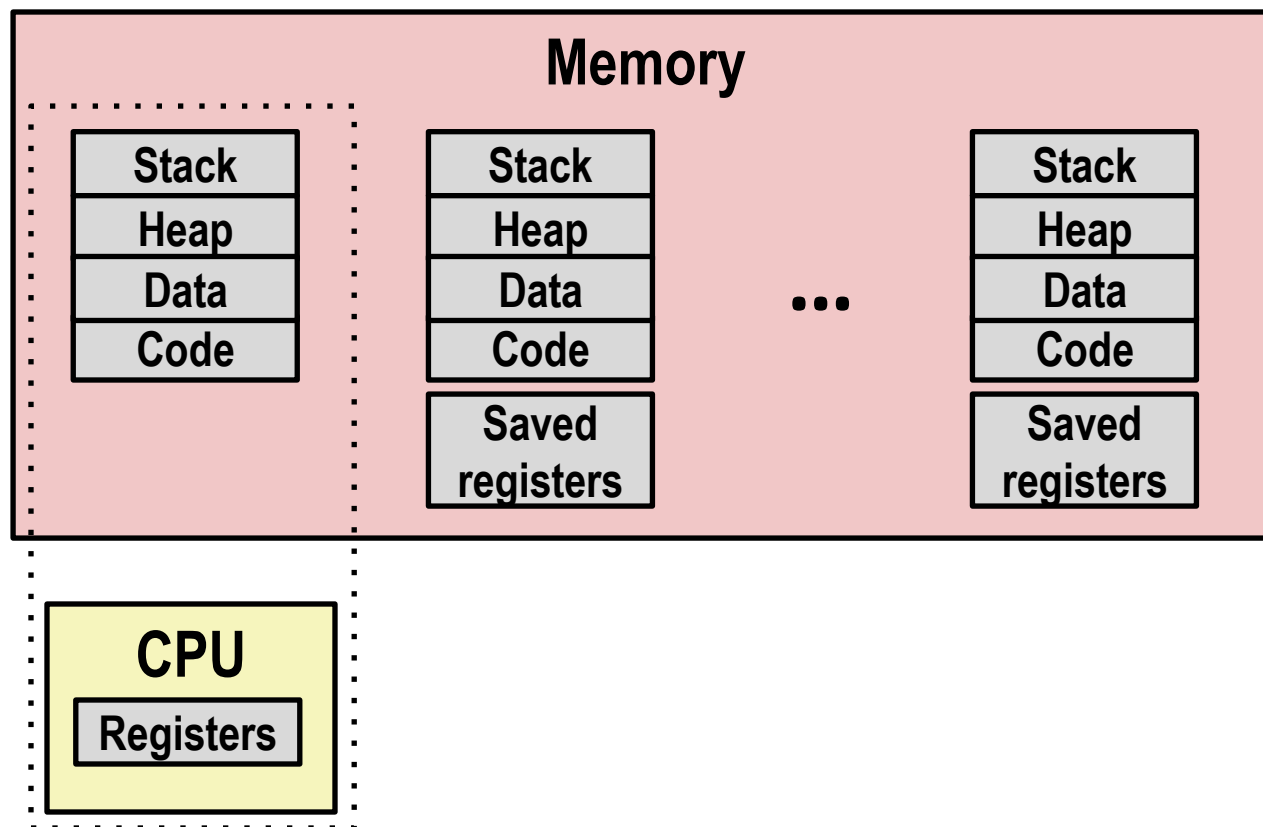
- Running program “top” on Mac or Linux
  - System has 123 processes, 5 of which are active
  - Identified by Process ID (PID)

# Context Switching

- Processes are managed by a shared chunk of memory-resident OS code called the *kernel*
  - Important: the kernel is not a separate process, but rather runs as part of some existing process.
- Control flow passes from one process to another via a *context switch*

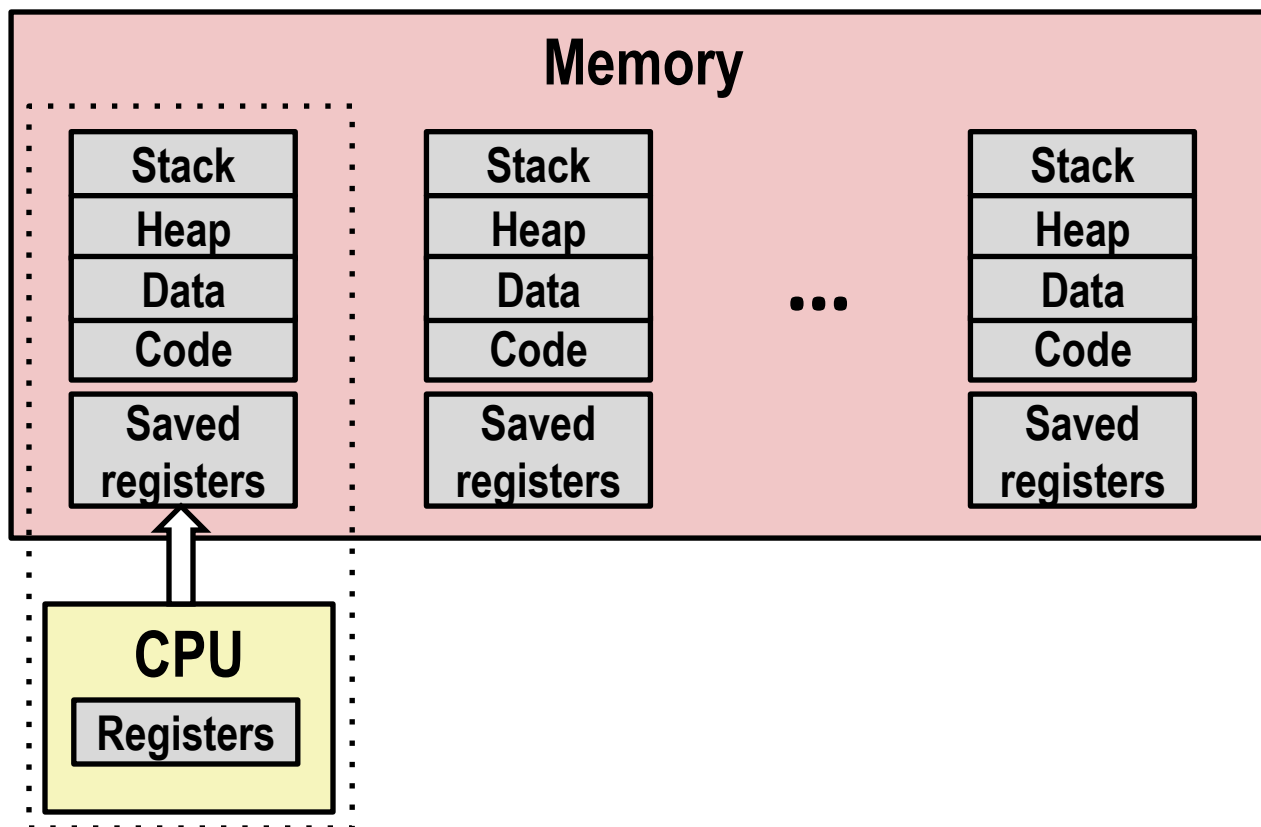


# Context Switching (Uniprocessor)



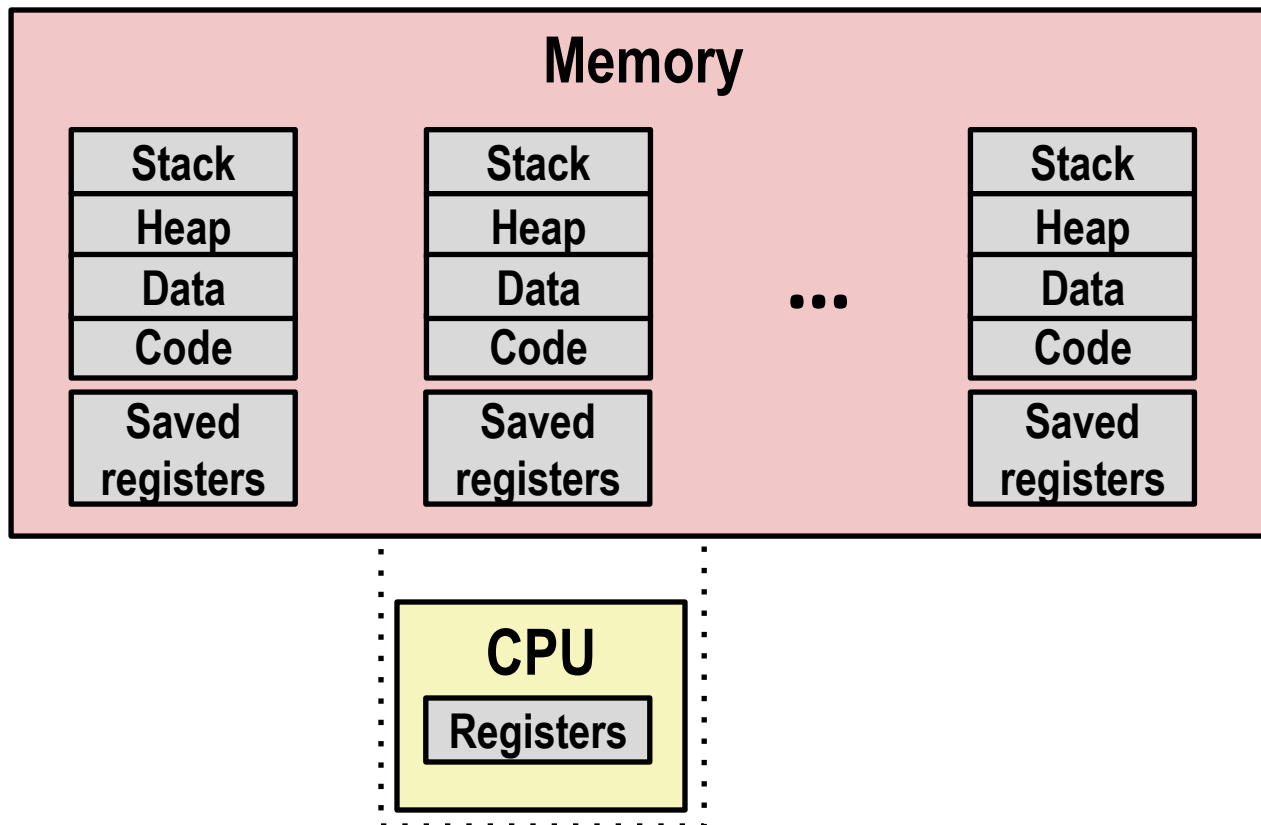
- **Single processor executes multiple processes concurrently**
  - Process executions interleaved (multitasking)
  - Address spaces managed by virtual memory system (like last week)
  - Register values for nonexecuting processes saved in memory

# Context Switching (Uniprocessor)



- **Save current registers in memory**

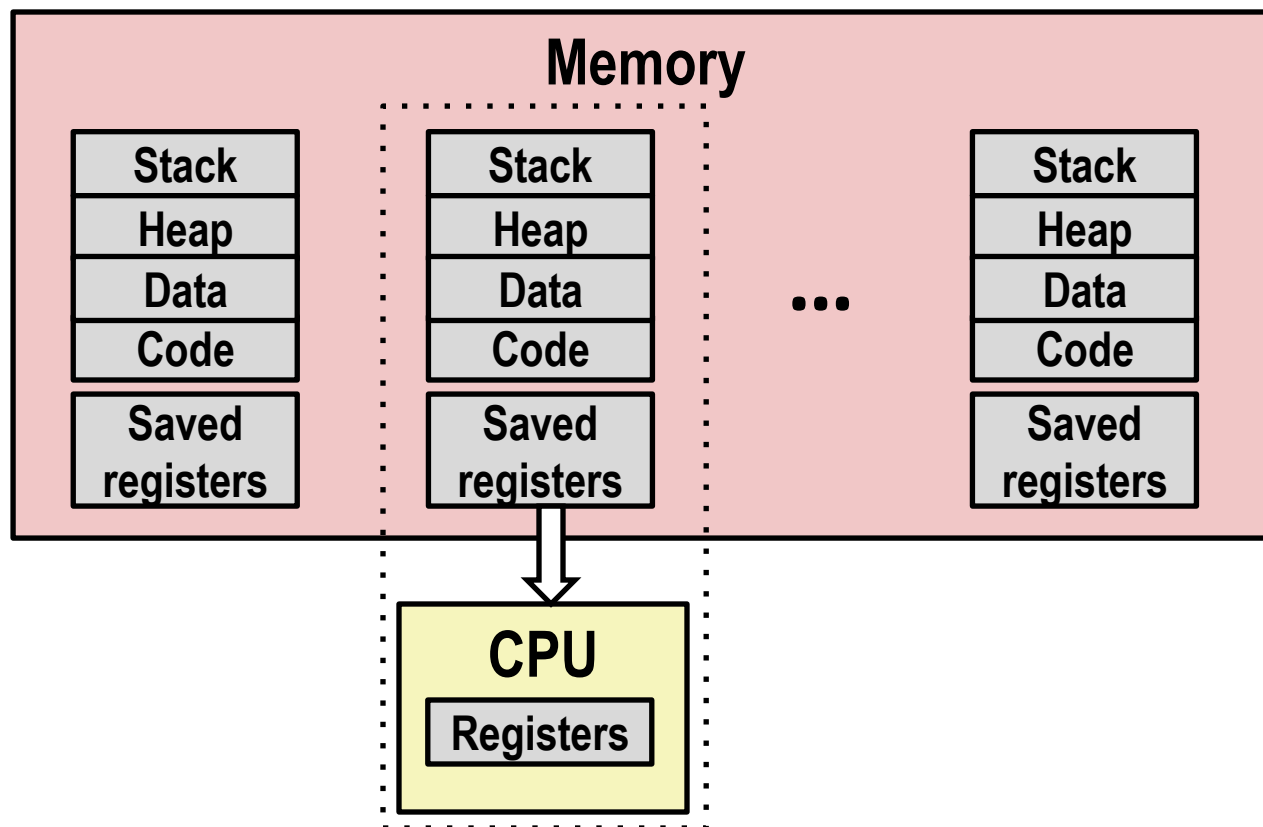
# Context Switching (Uniprocessor)



- **Schedule next process for execution**

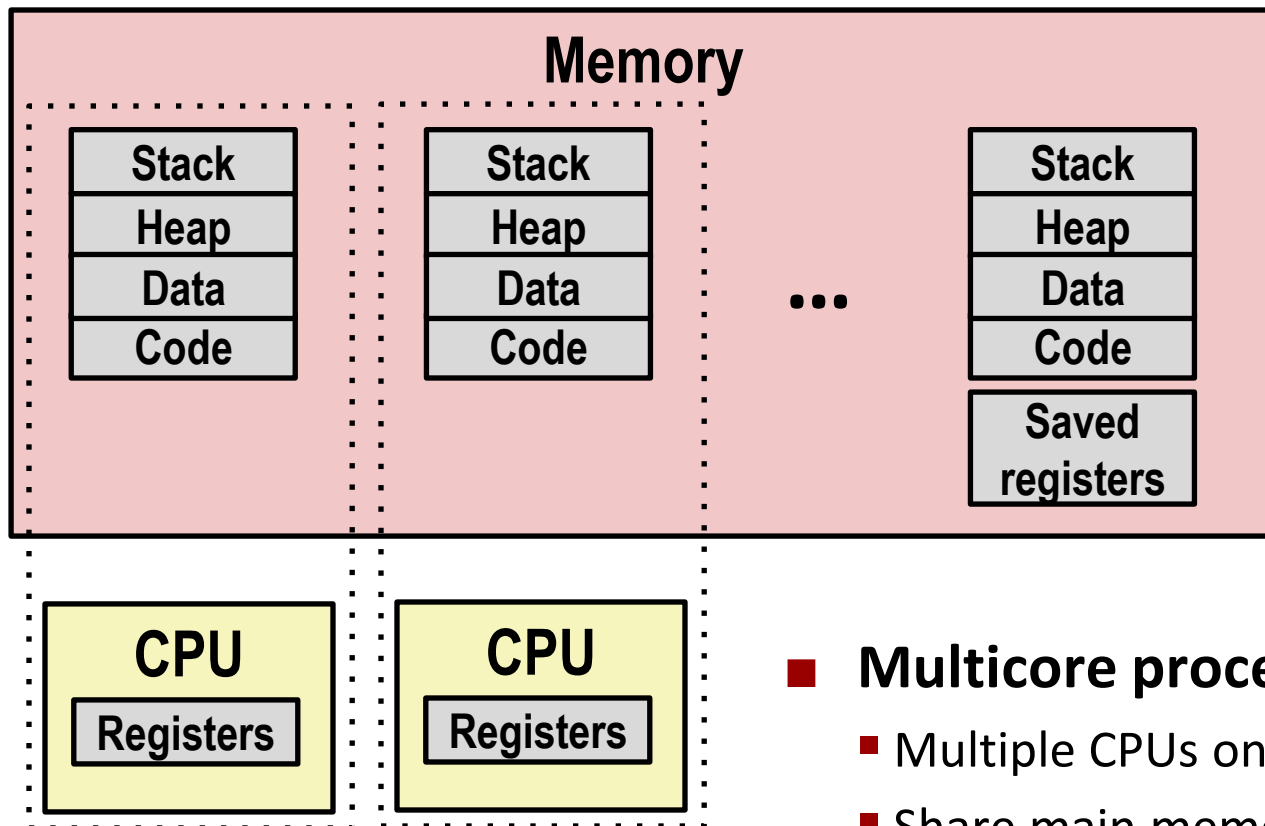


# Context Switching (Uniprocessor)



- Load saved registers and switch address space (context switch)

# Context Switching (Multicore)

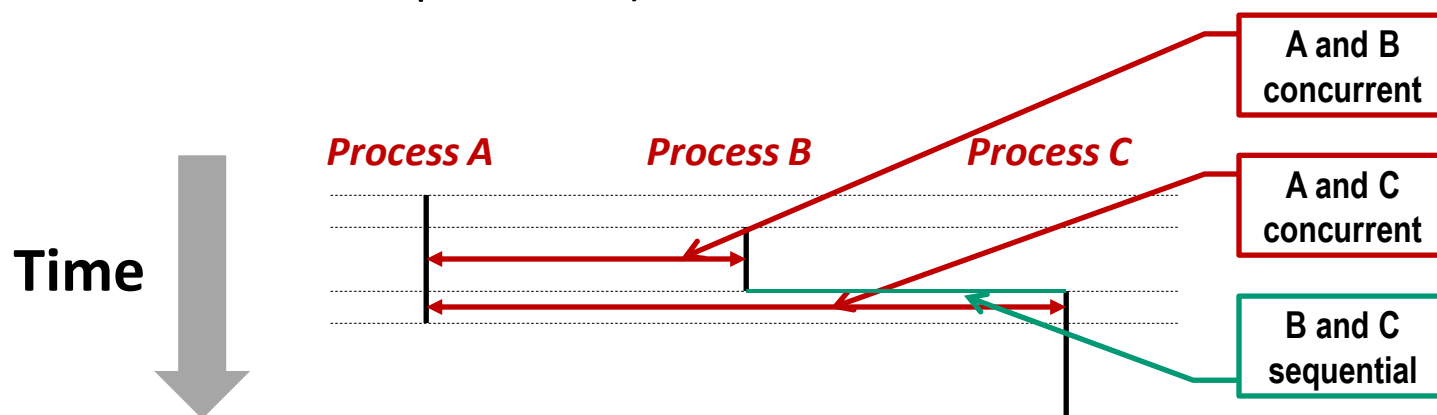


## ■ Multicore processors

- Multiple CPUs on single chip
- Share main memory (and some caches)
- Each can execute a separate process
  - Scheduling of processors onto cores done by kernel

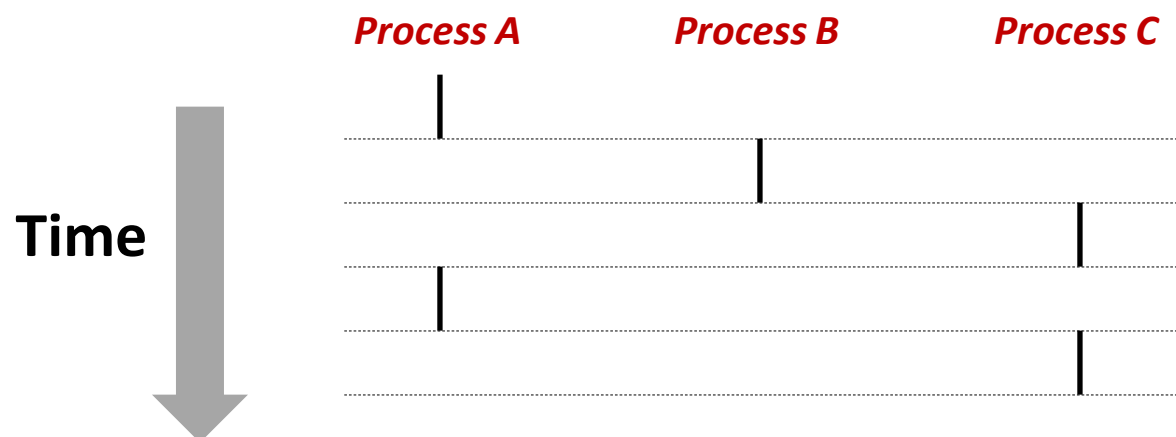
# User View of Concurrent Processes

- Two processes *run concurrently* (are concurrent) if their execution overlaps in time
- Otherwise, they are *sequential*
- Appears as if concurrent processes run in parallel with each other
  - This means they can interfere with each other (more on that in a couple weeks)



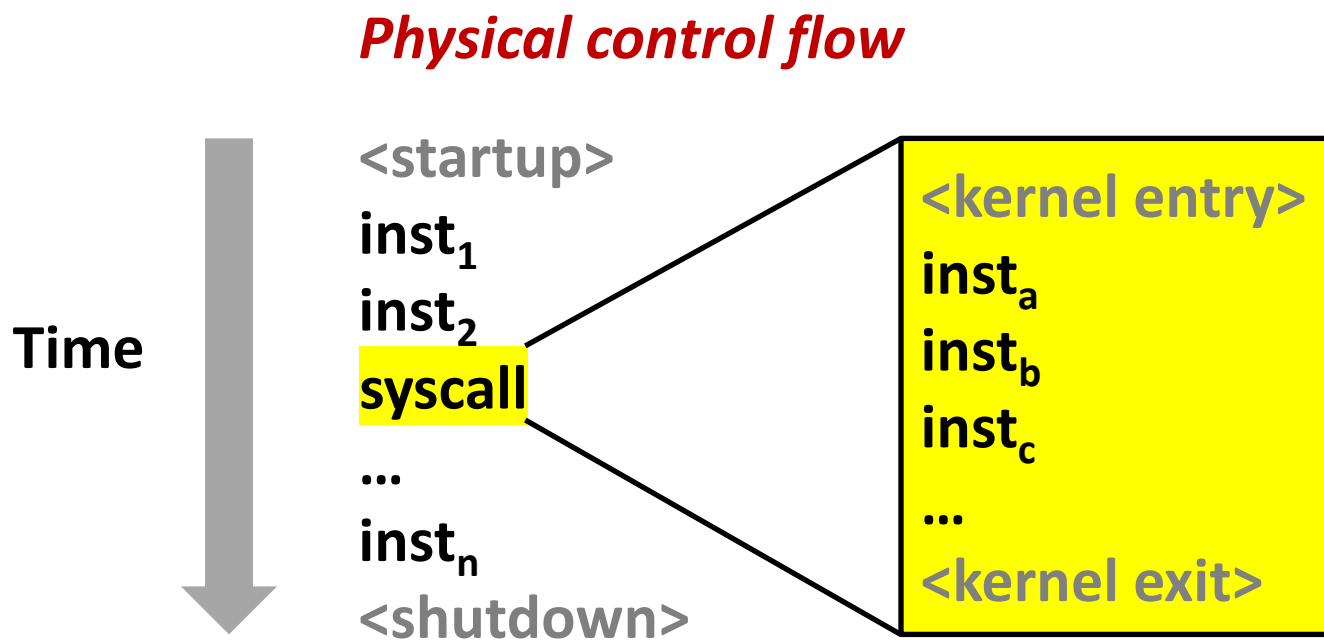
# Traditional (Uniprocessor) Reality

- Only one process runs at a time
- A and B execution is *interleaved*, not truly concurrent
- Similarly for A and C
- Still possible for A and B / A and C to interfere with each other



# How does the kernel take control?

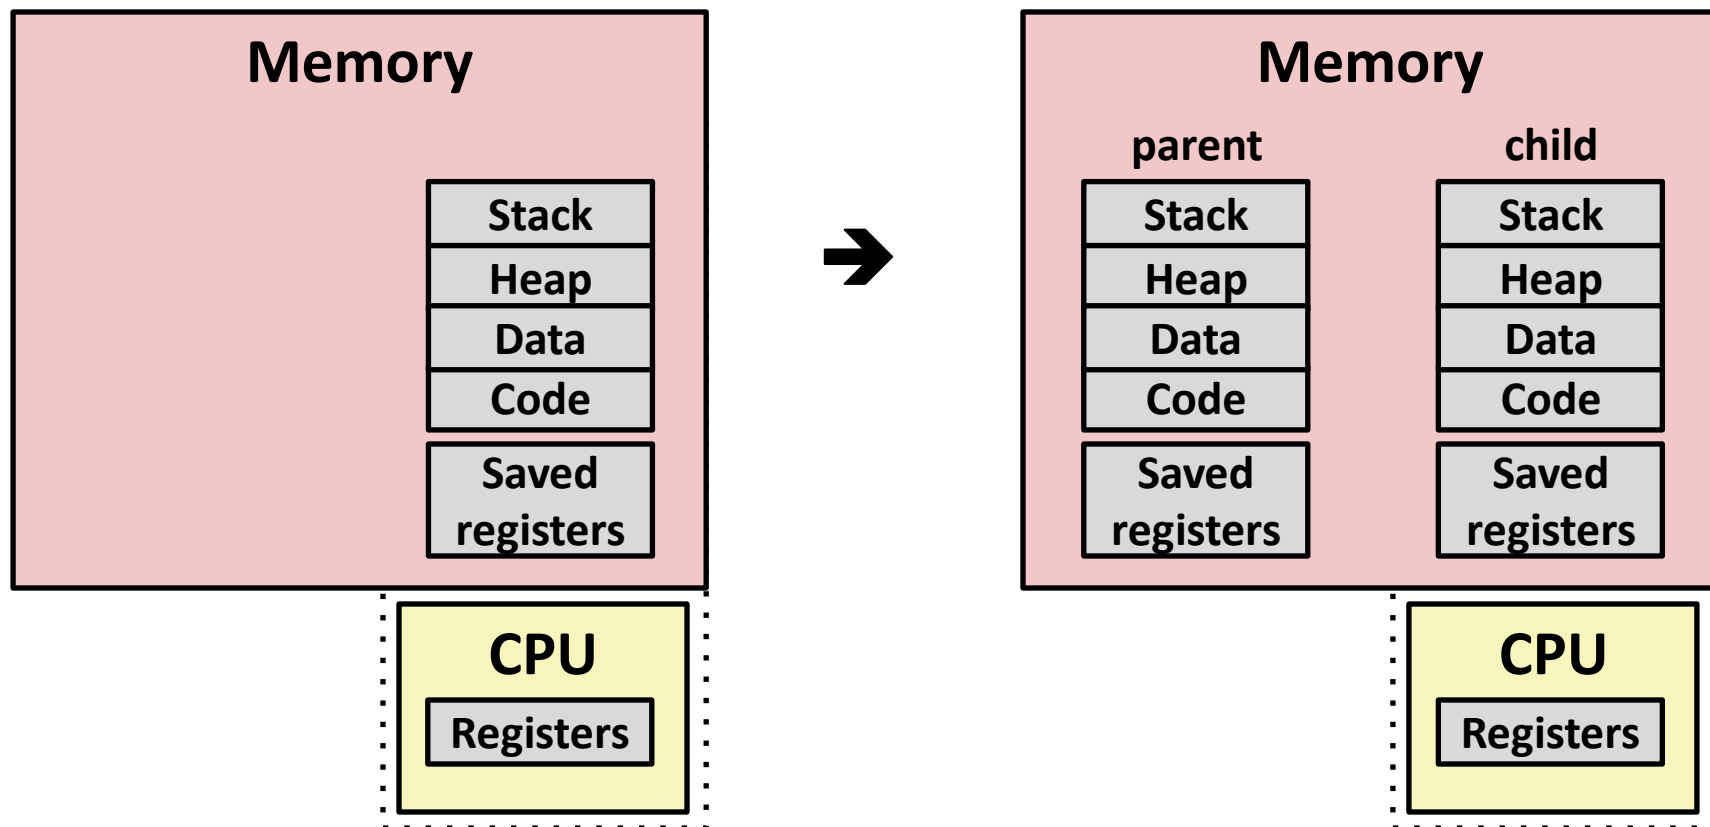
- The CPU executes instructions in sequence
- We don't write "now run kernel code" in our programs...
  - *Or do we??*



# Creating Processes

- *Parent process* creates a new running *child process* by calling `fork`
- `int fork(void)`
  - Returns 0 to the child process, child's PID to parent process
  - Child is *almost* identical to parent:
    - Child get an identical (but separate) copy of the parent's virtual address space.
    - Child gets identical copies of the parent's open file descriptors
    - Child has a different PID than the parent
- `fork` is interesting (and often confusing) because it is called *once* but returns *twice*

# Conceptual View of fork



## ■ Make complete copy of execution state

- Designate one as parent and one as child
- Resume execution of parent **or** child

# fork Example

```
int main(int argc, char** argv)
{
    pid_t pid;
    int x = 1;

    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        return 0;
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    return 0;
}
```

*fork.c*

- Call once, return twice
- Concurrent execution
  - Can't predict execution order of parent and child

```
linux> ./fork
parent: x=0
child : x=2
```

```
linux> ./fork
child : x=2
parent: x=0
```

```
linux> ./fork
parent: x=0
child : x=2
```

```
linux> ./fork
parent: x=0
child : x=2
```



# fork Example

```
int main(int argc, char** argv)
{
    pid_t pid;
    int x = 1;

    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        return 0;
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    return 0;
}
```

```
linux> ./fork
parent: x=0
child : x=2
```

- Call once, return twice
- Concurrent execution
  - Can't predict execution order of parent and child
- Duplicate but separate address space
  - `x` has a value of 1 when fork returns in parent and child
  - Subsequent changes to `x` are independent
- Shared open files
  - `stdout` is the same in both parent and child

# Summary

## ■ Exceptions

- Events that require nonstandard control flow
- Generated externally (interrupts) or internally (traps and faults)

## ■ Processes

- At any given time, system has multiple active processes
- Only one can execute at a time on any single core
- Each process appears to have total control of processor + private memory space

# 教材阅读

- 第8章 8.1.1-8.1.2、 8.2